

AINT351 - Revision for Ian

Three types of learning

Imagine a machine experiences a sequence of sensory inputs x_1, x_2, \dots, x_n

Supervised learning:

- The machine is also given y_1, y_2, \dots, y_n and its goal is to learn and reproduce them from the inputs
- Learning by examples, input and output is given so it knows how to reproduce the output from the input

Unsupervised learning:

- The machine should build a representation of x that can be used for decision making, prediction
- There is no desired output, you are given inputs and after some iterations you start to categorise data based on some criteria

Reinforcement learning:

- The machine can generate actions a_1, a_2, \dots, a_n that affects its environment and receives a reward or punishment based on them. Its goal is to learn actions that maximise long term reward
- Learning based on rewards for actions so that it learns to maximise long term reward

Goals of supervised learning

Classify input data:

- In this case the desired outputs y_1, y_2, \dots, y_n are discrete class labels and the goal is to **classify** new output correctly from the new input
- have an image of a digit and want to know what digit it is based on previous examples of that digit

Goals of unsupervised learning

Regression

- In this case the desired outputs y_1, y_2, \dots, y_n are continuous values and the goal is to **predict** new output correctly from new input
- Have the data from babies and can try to predict its weight given its height

We wish to find useful representations of data. This can involve

- Finding clusters
- Dimensionality reduction

- Finding the hidden cause of the surface phenomena
- Modelling the data probability density
- Data compression

Probability

Types of data:

Discrete data: only certain values

- Dice value = $\{1,2,3,4,5,6\}$
- Flip a coin = $\{H,T\}$

Continuous data: any value

- Length measurement
- Weight measurement

Probability functions

- A probability function maps possible values of a variable to its respective probabilities
 - e.g. if value is x we can write its possible probabilities as $p(x)$
- Probability functions have the following properties
 - $P(x)$ is a number with a value between 0 to 1.0
 - The area under a probability function is always unity

The addition law of probability

- If two events A and B are mutually exclusive then
- $P(A \cup B)$ = the probability event A **OR** B occurs
- $P(A \cup B) = P(A) + P(B)$
- If two events A and B are **NOT** mutually exclusive then
- $P(A \cup B) = P(A) + P(B) - P(A \cap B)$
- You have to subtract the intersect as it is where both events happen

Probability distributions

Bernoulli distribution:

- The probability of a success or failure, heads or tails, 1 or 0
- n is the number of times that the experiment is repeated

Discrete distribution:

- A finite amount of probabilities all of which have equal probability of occurring
- A dice throw, each outcome has a probability = $1/6$

Cumulative probability:

- The probability of this event happening **AS WELL AS** all the previous events
- A dice landing on 6 as well as all the chances of it landing on 1,2,3,4 and 5 = $6/6$

Binomial distribution:

- 2 outcomes
 - Heads or tails
- What is the probability of getting exactly 3 heads in 5 coin tosses
- HHHTT
 - $(1/2)^3 \times (1/2)^2$
- THHHT
 - $(1/2)^1 \times (1/2)^3 \times (1/2)^1$
- All equal = $(1/2)^3 \times (1/2)^2$
- therefore the overall probability =
 - $N \times (1/2)^3 \times (1/2)^2$
 - where N = number of unique arrangements
- There are exactly 10 ways to get 3 heads in 5 coin tosses
 - $N = 10$
- $10 \times (1/2)^3 \times (1/2)^2 = 0.3125$

Uniform distribution:

- A distribution that has constant probability
- 0.5 of values 0 and 1.0

Continuous data distributions:

- A continuous random variable is a random variable with a set of possible values that is infinite or uncountable
- looks like Gaussian distribution

Variance

- It measures how far a set of random numbers are spread out from their mean
- variance is the expectation of the squared deviation of a random variable from its mean

Expected value

- What is the expected value of x given it is in a certain area under a curve
- To find:
 - Get the marginal distribution by summing all the probabilities in that row or column (for that value of x or y)
 - Multiply the value of x by the marginal distribution of that value

Covariance: Joint probability

- The covariance measures the strength of the linear relationship between two variables
- to get the covariance of expected values:
 - find the expected value of x and expected value of y by using steps described above
 - get the expected value of $xy = E(XY)$ by multiplying each value in the table by its x and y values
 - and then plug all values into equation
 - $COV(XY) = E(XY) - E(X) \cdot E(Y)$

Coefficient of XY

- To get the coefficient of XY we need the standard deviation of x and standard deviation of y as well as the covariance figured out in steps above
- $Coff(XY) = COV(XY) / STD(X) \cdot STD(Y)$
- First get the variance of x and y
 - **variance of $x = E(X^2) - E(X)^2$** - to get the $E(X^2)$ need to do it the same as the expected values by square all the values of x
 - **variance of $y = E(Y^2) - E(Y)^2$** - to get the $E(Y^2)$ need to do it the same as the expected values by square all the values of Y
- Once you have the variance you can square root it to get the standard deviation
- then plug it all back into the equation

Conditional probability distribution, independence

- How can you tell if x and y are independent
- Changing the value of y should have no effect on the probability distribution of x

Effect of standard deviation

- The variance either side of the mean

- The greater the standard deviation the greater the probability that x is within it
- Greater the standard deviation the further it is from the mean

Cumulative distribution function

- For a continuous random variable X the cumulative distribution function
- CDF(x) represents the area under the probability density function P(x) to the left of X
- $CFD(x) = P(X < x)$

Exponential distribution

- Given $P(x) = e^{-x}$
- What is the probability of x falling withing 1 to 2
- $P(1 \leq x \leq 2)$ corresponds to area under distribution between 1 and 2

Marginalization

- sum up all the columns and row values for either x or y respectively

Conditional probability: bayes rule

- 'P(A n B) is the probability of A and B happening
- can be written P(A, B) - the joint distribution of A and B
- The probability of A happening multiplied by the probability of B happening given that A has happened

$$P(A \cap B) = P(B|A)P(A)$$

$$P(B|A) = \frac{P(A \cap B)}{P(A)}$$

- if A and B are **NOT independent events** then - $P(A \cap B) = P(A) \cdot P(A|B)$
- **This relationship is always true**

Interpreting covariance

- If we calculate the covariance between two random variables
- if the $\text{cov}(X, Y) > 0$
 - X and Y are positively correlated
- If the $\text{cov}(X, Y) < 0$
 - X and Y are inversely correlated
- If the $\text{cov}(X, Y) = 0$
 - X and Y are independent

Marginalisation

- Sum of probabilities across their given variable

Conditional probability

- If A and B are **NOT** independent events then
- $P(A \cap B) = P(A) \cdot P(B|A)$
- This relationship is always true
- A has to happen for it to be true
- If A and B **ARE** independent events then
- $P(A \cap B) = P(A) \cdot P(B)$
- This relationship is only true if A and B are independent
- Rolling a dice:
 - The first throw doesn't effect the second

Product rule of probability

- Product rule states that
- $P(A, B) = P(B)P(A|B)$
- $P(A, B) = P(A)P(B|A)$
 - The joint probability of A and B is prob of A multiplied by the probability of A given B
 - * Same the other way around
- Leads to bayes rule
- $P(B)P(A|B) = P(A)P(B|A)$
- $P(A|B) = P(A)P(B|A)/P(B)$
 - The probability of A given B is equal to:
 - The probability of A multiplied by the probability of B given A


- Over the probability of B
 - This is bayes theorem
- If A and B are independent
 - $P(A, B) = P(A)P(B)$
 - Joint probability of two independent events is the dot product probability of both events happening

Gradient descent

- How can we get to the top of a hill?
 - We follow the gradient
- How can we get to the bottom of a hill
 - We follow the descent
- How do we know we're at the top
 - It goes down on both sides

Gradient of a curve

- Gradient is the slope of a curve or surface
- Going up the hill it is +ve
- Going downhill it is -ve
- Differentiation finds tangent of line on graph
- Gradient of straight line =
 - Change in y
 - over change in x



Iterative gradient decent

- Find minimum of a function

- Move downwards in direction of gradient

Local maxima and minima

- Can be local minima as well as global minima
 - Same for maxima
- Gradient descent can get stuck in local minima/maxima

Least squares fitting

- We want to fit a straight line to data measurements
- Equation for straight line in a single dimension is:
 - $y_i = mx_i + c$
- Sum error over all points

Generate a N dimension Gaussian distribution

- If we use `randn(N, 1)` to draw N samples from a 1D distribution x_1
- And use it again to draw N samples from a 1D distribution x_2
- From these we build a 2D vector X
- What is the covariance matrix of 2D dataset X
 - $E = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$
- Covariance terms = 0 and 0
- Means that they are independant

Terminology for types of learning

Maximum likelihood (MP) learning:

- Does not assume a prior over the model parameters
- Finds q parameter settings that maximises the likelihood of the data
- ML may run into estimation problems if we're unlucky

Maximum a posteriori (MAP) learning:

- Assumes a prior over the model parameters
- Finds a parameter settings that maximises the posterior
- Having a prior belief means that the intial estimate may be more appropriate

Bayesian learning:

- Assumes a prior over the model parameters
- Computes the posterior of the parameters

NB: MAP and ML estimates are identical when the prior is uniformly distributed

Gaussian class-conditional model

- The maximum likelihood fit of a Gaussian to some data is the Gaussian whose mean is equal to the data mean and whose covariance is equal to the sample covariance
- One nice feature of this model is that the MP parameters can be found in closed-form:
 - So we don't have to use iterative solutions

Limitations:

- We cannot account for higher order statistical structure in the data
 - These require nonlinear and hierarchical models
- We need to deal with outliers
 - These require nonlinear and hierarchical models
- The multivariate model uses $D(D+1)/2$ parameters
 - If D is very large we need to use dimensionality reduction

Frequentest and Bayesian statistics

Frequentest approach:

- Probability is the limit of observed frequency as number of observations goes to infinity
- Considers the model parameters to be fixed (but unknown) and calculates the probability of the data given those parameters

Bayesian approach:

- Probability is a “degree of confidence” that one attaches to an uncertain event
- Requires a prior estimation of the models likelihood
- naturally incorporating prior knowledge

Eigenvalues and Eigenvectors

- If A is an $n \times n$ matrix
- A scalar λ is called an eigenvalue of A
 - If there is a non-zero vector x such that
 - $Ax = \lambda x$
- If we have a matrix and a scalar value
 - A and λ
- If both of them multiplied the vector x
 - result in the same vector values of x

- Then such a vector x is called an eigenvector of A corresponding to λ
- Eigenvectors can only be found for square matrices
- Not every square matrix has an eigenvector
- An $n \times n$ matrix can have n eigenvectors
- All the eigenvectors of a matrix are orthogonal
- The length of an eigenvector is exactly one?
- When we transform X by multiplying by A we end up with vector X again but this time scaled by λ
 - Therefore the **direction of vector X is unaffected by the transformation**
- Only keeps the vectors of parallel values

Clustering

- Idea of clustering is to group patterns so that:
 - Patterns that are similar to each other are in the same cluster
 - Patterns that are dissimilar to those are in the other clusters
- All require a way to determine similarity

K-means clustering

- K is the number of clusters **Steps:**
- Assign K
- Randomly assign the first K data points to be the centroids of the K clusters
- While loop
 - For each point, assign it to its closest cluster centre
 - Re-compute the cluster centres as the means of assigned data points
 - Terminate loop if there was change in assignment compared to last iteration (converges)
- K-means is an example of hard clustering
 - Clusters do not overlap
 - Data in one cluster only

Mixture of Gaussians

- Gaussian, FA and PCA models are easy to understand and use in practice
- They are a convenient way to reduce dimensionality of high dimensional data sets
- The problem is that they make very strong assumptions about the distribution of the data
 - Only the mean and variance of the data are taken into account
- The class of densities which can be modelled is also too restrictive
- By using mixtures of simple distributions
 - Such as Gaussians
 - We can expand the class of densities greatly
- Mixture of Gaussians is an example of soft clustering
 - Clusters may overlap
 - Data may exhibit non-binary strength of association to all clusters
- Mixture of Gaussians
 - Probabilistic method
 - Each cluster is a generative model
 - Clusters have parameters (mean & covariance)
 - Train using EM
- Advantages is that given enough components we can model most distributions using only the simple Gaussian distribution

EM Algorithm

- Need to know the Gaussian parameters (mean and standard deviation) for each cluster to estimate the data point cluster membership
- But need to know data assignments to estimate the parameters
- Solution is to use the EM algorithm
 - Randomly initialise the Gaussian cluster parameters
- E step:
 - Look at each data point and calculate how likely it came from a given cluster
 - Do so by computing $p(b|x_i)$
- M step:
 - Re-estimate Gaussian parameters to fit the assigned points
 - Repeat until convergence

Pattern classification

- Pattern classifiers partition the input space

- May have multiple input data dimensions
- May have multiple output classes
- Type of decision boundary depends on the classifier
- Variety of ways to determine boundaries

Naive Bayes classifiers

- Estimation of full high dimensionality covariance matrix can be a problem
 - Why?

Problem:

- We can only estimate $P(x_1, x_2, \dots, x_n | y)$ based on counts in the training dataset
- We may not see every x value for every y value

Solution:

- Make an assumption of independence between the features
- Assume conditional independence given y
- so for a dataset:
 - Calculate probabilities of each feature x_i independently
 - Then joint probability is just the product of individual feature probabilities

Conditional independence

Example - Going to the beach and getting heatstroke not independent - $P(B, S) > P(B) \cdot P(S)$ - However they may be independent if we know that the weather is hot - $P(B, S | H) = P(B|H) \cdot P(S|H)$ - In this case the hot weather explains a reason for dependence between going to the beach and getting heatstroke - Similarly in classification the class values explains the dependencies between the attributes

Information theory

- The probability $P(X)$ encodes uncertainty about the random variable X

Entropy

- We can quantify the information represented in such a random variable
- This is the entropy of the variable X which is average amount of information required to encode X

Mutual information

- It indicates how much knowledge if Y affects our belief in X
- If Y has no effect and is independent of X then the mutual information is zero

Naive Bayes versus full Gaussian classifier

- The probability of a data point
 - how probable is x_i under source c
- Naive bayes only estimates and uses marginal Gaussian parameters
- Only has non-zero values in the covariance matrix along its leading diagonal
- Assumes variables are independent
 - zero covariance
- Can only represent distributions aligned with coordinate system
- Easy to estimate parameters
- Choose most probable class given observation

Limitations of generative models

- Bayes' decision rule minimise average probability of error
- Can train generative models by directly estimating parameters
- So what needs improving
- Bayes classifier is the minimum error classifier only if our model of the data is appropriate
- In particular if the form of the class conditional distribution is correct

Discriminative models

- Generate model classification requires the class posterior
- Discriminative models use parameters more closely related to classification process
- They are not dependent on generative process being correct
- Don't model the data just try to find decision boundary
- Example is the perceptron

Linear decision boundaries

- Decision boundaries partition the input space into discrete regions
- Each region is associated with a class label

The perceptron

- Learning means changing the weights between the neurons
- Relationship between input and output is important in computational neuroscience
- Simple but limited capabilities
- Basic concepts are useful for multi-layer models
- If the data is linearly separable then:
 - A linear decision boundary will correctly classify all points
 - Algorithm will stop where there are no incorrectly classified training data points
 - The algorithm therefore guaranteed to converge to a solution
- Two forms of update are generally adopted
- Batch update:
 - All the data is presented for each iteration
 - Weights only updated after all data seen
- Sequential update
 - Data is presented one sample at a time
 - Weights updated after each sample

What is deep learning?

- ‘*Deep learning*’ involves using a neural network with several layers of nodes between the input and output
- A series of layers between the input and output do feature detection and processing in a series of stages
- Model of human visual system
- Use recent algorithms for training many-layer networks

Deep networks

- Two main types of deep networks:
- Convolutional neural networks
 - Employ alternating layers of convolutional networks followed by a pooling layer
 - Output uses traditional MLP
- Deep belief networks
 - Consist of perceptron stacked Boltzmann machines
 - Use classification output layer

Auto encoder

- An auto-encoder is trained with standard training algorithm

- It learns to map the input back onto itself

New way to train multi-layer NNs

- Use greedy layer-wise training to train multilayer networks
- Train first layer using data unsupervised without labels
- Use abundant unlabeled data which is not part of the training set
- Freeze the first layer parameters and start training the second layer using the output of the first layer as the unsupervised input to the second layer
- Repeat this for as many layers as desired
- this builds a set of robust features
- Use the outputs of the final layer as inputs to a supervised layer/model and train the last supervised layers
 - leave early weights frozen
- Each of the non-output layers is trained to be an auto-encoder
- Essentially it is forced to learn good features that describe what comes from the previous layer

Advantages of greedy layer-wise training

- Avoids many of the problems of trying to train a deep net in a supervised fashion
- Each layer gets full learning focus in its turn since it is the only current ‘top’ layer
- Can take advantage of unlabelled data
- When you finally tune the entire network with supervised training the network weights have already been adjusted so that you are in a good error basin and just need fine tuning
- This helps with problems of
 - Ineffective early layer learning
 - Deep network local minima

Convolution networks

The replicated feature approach

- Use many different copies of the same feature detector with different positions
 - Could also replicate across scale and orientation
 - Replication greatly reduces the number of free parameters to be learned
- Use several different feature types, each with its own map of replicated detectors
 - Allows each patch of image to be represented in several ways

Effect of replicating feature detectors

- **Equivariant activities:** Replicated features do **not** make the neural activities invariant to translation
 - The activities are equivariant
- **Invariant knowledge:** If a feature is useful in some locations during training
 - detectors for that feature will be available in all locations during testing

Pooling replicated feature detectors outputs

- Get a small amount of translational invariance at each level by averaging four neighbouring replicated detectors to give a single output to the next level
- This reduces the number of inputs to the next layer of feature extraction
 - Thus allowing us to have many more different feature maps
- Taking the maximum of the four works slightly better
- **Problem**
 - After several levels of pooling we have lost information about the precise positions of things
 - This makes it impossible to use the precise spatial relationships between high-level parts for recognition

Convolutional networks

- Convolution is the integral of the product of the two functions after one is reversed and shifted
- Same as sliding a single instance of the detector over input
- Performs convolution like an FIR filter
- Each convolution layer contains multiple feature maps
 - Just as hidden layers typically multiple nodes
 - So each map is a single feature detector
- Repeatedly applied to a small window across the whole input
- A convolutional neural network is tiled in such a way that they respond to overlapping regions in the input field
- Convolutional layer consists of units that act as feature detectors
- Each convolutional unit is connected to a small region of the input
- This structure is replicated for all offsets across the input of data to span the input space

- Thus the augmented connection weights are shared for all the hidden units in this feature map
- Each layer combines patches from previous layers
- Typically tries to compress large data into smaller set of robust features based on local variations
- Convolution can create many features

Pooling layer

- A pooling layer often follows a convolution layer
- It down samples the feature map
- This step can compress and smooth the data
- Make data invariant to small translational changes
- Usually takes the average or max value across disjoint patches
- Feature outputs are pooled independently

Final MLP layer

- To make decisions a final MLP is often used
- This can then make discrete classifications of the input

The perceptron

- Bias can be augmented into weight vector
- Also add corresponding unity value in augmented data vector

The perceptron learning rule

- TO train the perceptron
- first initially set all weights to small random values
- Then loop
 - For all data point and its output target in turn, compute perceptron output z (either 1 or a 0)
 - After each data point presentation the weight update is given
 - Weight of input i perceptron node
 - multiplied by the target + 1 is equal to
 - this weights input i perceptron node
 - multiplied by the target
 - plus the learning rate multiplied by target for current input
 - minus the current output
 - times x which is the i th input

- Need to repeat procedure across all the dataset until finished
- Guaranteed to converge if the problem is separable

Perceptron rule versus delta rule

- Perceptron provides learning feedback after a threshold non-linearity
- Instead can provide feedback before the threshold
- Replace step function with continuous differentiable function
- Can be linear - if so, problem similar that of linear regression
- can still use threshold for final classification
- This has the advantage we can use gradient descent to find weights

Single vs multi layer networks

- Single layer networks implements linear decision boundary
- Multilayer networks can implement more complex decision boundaries

Kernel trick

- Transforming some datasets can make it linear separable
- more rigorously this can be done using a kernel
- A kernel is basically a mapping function that transforms one given space into another
- Such transformation of data that leads to easier to solve problems is sometimes known as the kernel trick

Limitation of back propagation

- Multiple hidden layers
- get stuck in local optima
 - Start weights from random positions
- Slow convergence to optimum
 - Large training set needed
- Only use labelled data
 - Most data is unlabelled
- Error attenuation with deep nets
- Can now training with GPUS and special hardware
- However can still be time intensive for large networks