

SOFT354 - Increasing performance of CUDA Programs

Date: 13-10-16

2D Blocks and Grids

Remember: - A **block** consists of multiple **threads** - a **grid** consists of multiple **blocks**

- Number of blocks to use depends on what you're trying to process
- Use of experimentation to get what is right for you
 - Make code function without dependency on how many blocks it uses
- Use dim3 objects to launch kernel with grid/block with dimensions greater than one
- To launch a kernel with grid / block dimensions that aren't one dimensional use **dim3** objects.

```
dim3 gridSize(94, 125, 1);  
dim3 blockSize(32, 32, 1);  
invertKernel<<<gridSize, blockSize>>>();
```

Need to give z
dimensions as well
- just set to 1.

- Kernel can now determine its **global position** in the image using **blockIdx**, **blockDim** and **threadIdx**:

```
__global__ void invertKernel() {  
    int globalIdx = blockIdx.x * blockDim.x + threadIdx.x;  
    int globalIdY = blockIdx.y * blockDim.y + threadIdx.y;  
    if (globalIdx < 2992 && globalIdY < 4000) {  
        // Do something...  
    }  
}
```

- Use globalId of x and y to get the position of the threads calculation within the grid as a whole - Clause in if statement checking the global id gets rid of threads that don't have a thing to process if you had to round up on number of blocks

- Can do the same but include z direction with 3D blocks
- Z dimensional isn't as high as X or Y
 - $X/Y = 1024$
 - $Z = 64$

```
dim3 gridSize(100, 100, 100);
dim3 blockSize(10, 10, 10);
someKernel<<<gridSize, blockSize>>>();
```

```
__global__ void someKernel() {
    int globalIdx = blockIdx.x * blockDim.x + threadIdx.x;
    int globalIdxY = blockIdx.y * blockDim.y + threadIdx.y;
    int globalIdxZ = blockIdx.z * blockDim.z + threadIdx.z;
```

Linearisation of blocks and grids

- 2D and 3D blocks and grids are just for convenience
- CUDA will convert you **blocks** into a 1D array of **threads**
 - .. and your **grid** into a 1D array of **blocks**

Block assignment

- A GPU has multiple *streaming multiprocessors* (SMs)
- Each SM has one or more blocks assigned to it (“**resident**”) at any given time
- Only **whole blocks** can be assigned to SMs, therefore **all threads in a block will run on the same SM**
 - Important for memory sharing

Warps

- When a streaming multiprocessor has had some blocks assigned to it
- when it takes the block, it breaks it down into warps
- **warps** are a set of 32 **threads**
 - These will be executed in parallel

Warp Scheduling

- Generally you don’t need to worry about warps
- But by understanding how an SM schedules warps you can improve performance
- Only one warp can be running at any given time
- Any given time a warp can be in one of three states:
 1. **Running** currently being executed by the SM
 2. **Waiting** Can’t run because it’s waiting on something - usually some data to read from memory
 3. **Ready** Not waiting but is ready to run once it is its turn

- Ideal situation is when there is one currently running, two are ready to run when it is finished
- Bad situation is if they're all waiting
- If all warps were waiting, none could be executed and therefore they are **stalled**
 - Usually waiting for memory access

Reducing stalls

- The most effective way to reduce stalls (assuming they're mostly caused by memory access) is to use memory more effectively
 - Coalesced memory reads and shared memory
- But it is also important to maximise the number of warps resident in an SM (its **occupancy**);
- The more resident warps, the more chance there's one **ready** to be executed

Occupancy

- When looking at how many warps can be resident in a SM
- Points to consider
 - Max no of threads per block
 - Max no of resident blocks per multiprocessor
 - Max no of resident warps per multiprocessor
 - Max no of resident threads per multiprocessor
- **8x8 blocks:**
 - 64 threads per block (2 warps).
 - But each SM can only contain up to 8 blocks.
 - So maximum number of resident threads is $8 \times 64 = 512$.
 - Much less than the capacity (1536 threads).
- **16x16 blocks:**
 - 256 threads per block (8 warps)
 - If SM contains 6 blocks, $6 \times 256 = 1536$ threads.
 - Maximum possible occupancy – very good!

- 32x32 blocks:

- 1024 threads per block (32 warps).
- Maximum capacity is 1536 threads.
- But can only allocate whole blocks to a SM, so there's only space for one.
- So maximum number of resident threads is 1024.
- Much less than the capacity (1536 threads).

- There are limits to how many registers and shared memory a SM has - This can also limit occupancy

Coalesced Memory access

- When memory being accessed is within 128 bytes of each other
 - You can access the all memory withing that range
- This saves optimistation as you don't have to multiple calls to the memory

Shared memory

- We've already met two types of statically allocated device memory:
 - **device** variables are stored in **global memory** which is *huge but slow*
 - **constant** vairables are cached in the **SM**; can be read by all threads *in a single cycle* but are *read-only*
- The last to consider is **shared**
 - How the data should be put together
- A variable with the **shared** modifier has its memory allocated in the SM's shared memory area
- This is **fast** but local to the SM
- Each block has a separate copy of any **shared** variables
- All threads in a block can access their block's copy of the shared variables, but not other blocks' copies
- Shared memory can be used to allow threads **within a block** to communicate with each other very quickly
- Common pattern is when all threads in a block need the same chunk of global memory:
 - Each thread loads a bit of the global data into shared memory (happens in parallel)
 - *Library metaphor*: Study group makes a list of books they need, everyone fetches a book from gloabl memory (shelves) and puts it on a shared desk

`__syncthreads()`

- Function that can only be called from the kernel code
- No guarantee if threads have done their job without `__syncthreads()`
- Causes all threads in the current block to wait, until they've all reached the call

Conclusions

- **Threads** are organised into **blocks**, which are organised into a **grid**
- The **threads** in a **block** and/or the **blocks** in a **grid** can be organised as 1D, 2D or 3D structures
- Whatever the dimensionality, **blocks** and **threads** are linearised into a big 1D array of threads
- Whole **blocks** are allocated to streaming multiprocessors as they become available
- When a block is allocated to a multiprocessor it is divided into **warps** of 32 threads
- Want to avoid the situation where all **warps** in a SM are waiting and not runnable, Two strategies:
 - Maximise the number of warps in an SM (its *occupancy*) by optimising block size
 - Reduce the time spent waiting for memory access by using **coalescing** and **shared memory**