

## SOFT354 - Revision notes

Date: 03-12-16

### Static allocation

- Size of array must be known at compile time
- Space is reserved in the program memory map

*Allocation in host memory:* `float h_array[10];`

*Allocation in device memory:* `__device__ float d_array[10];`

### Dynamic allocation

- Size of array can be calculated at run time

*Allocation in host memory:* `float* h_array = (float*)malloc(10*sizeof(float));`

*Allocation in device memory:* `float* d_array = (float*)cudaMalloc(&d_array, 10*sizeof(float));`

- If you use `malloc` inside the kernel it will allocate memory on the **device**
- If you use `cudaMalloc` inside the kernel
  - Every thread would allocate enough space to hold one copy of the array in device memory
  - And each thread accesses a different bit of it

### Copying to static device arrays

- When you use **static allocation** for device memory
  - `__device__ float d_array[10];`
- The variable `d_array` is not a pointer to a memory address on the GPU
  - **it is a symbol**
- Therefore you cannot use `cudaMemcpy`
- Have to use `cudaMemcpy*Symbol`

### Accessing static memory from a kernel

- If device memory was allocated statically in the host
- It allows it to be accessed globally on the device
- It doesn't require time consuming calls to memory allocation functions
- Automatically frees itself in memory after use

## Pinned Memory

- Technique for speeding up RAM < - > GPU memory transfers
- Transfers between RAM and GPU use **Direct Memory Access** to do the copying without working the CPU
- Where by a section in RAM is 'pinned' for the GPU as RAM uses paging to swap around memory
- This means that the pages aren't swapped so it allows for faster transactions as it doesn't require finding the memory address
- Because of paged virtual memory
- The address you get from `malloc()` won't correspond to a physical address in RAM but instead a **page** that can be moved around
- So it copies it first into DMA where it won't get moved
  - Then that allows for quicker transfer
- Transfers can happen entirely with this
  - big speed increase

## cudaMallocHost

- like `malloc` it dynamically allocates an array in RAM
- unlike `malloc`, the memory will be **pinned**
  - Memory calls to `cudaMemcpy` will be faster

## Uniform cache

- The uniform cache is an on-chip (part of each SM) cache that is designed for **broadcasting data**
- If multiple threads access the same address in the uniform cache at the same time
  - The data is sent to them all simultaneously
- Threads can't change the value of data in the uniform cache
- Similar to L1 cache
- SM has its own L1 and uniform cache
- It is used for broadcasting out data

## Registers

- Fastest, accessible by a single thread
- The fastest
- Local variable inside kernel go into registers where possible
- Two exceptions:
  1. If the index accessing the array using a variable then it is stored in local memory
  2. If you use more than the available register space

## L1 Cache / Shared memory

- located in an SM
- Shared between threads in that SM
- Extremely fast
- Individual to each SM
  - L2 is shared between all SMs
- **shared:** you have to program its use
  - Allocated by using `__shared__` modifier
- Can control how it is used
- Every block of thread has this area to communicate with threads that are in the same block
  - Each block has a separate copy of shared variables
  - All threads in a block can access their copy of variables but not other SMs
- Common pattern:
  - All threads in a block need the same chunk of global memory
  - Each thread loads a list of the global data into shared memory in parallel

## L2 Cache

- Shared between all SMs
- Much faster than RAM but slower than L1
- Bigger than L1

## Global memory - RAM

- Biggest and slowest
- Shared between all SMs
- **Only memory we can read and write from host**
- Only memory that you can directly access from inside the host - using `cudaMemcpy`
- Slowest type of transfer
- Can speed it up using DMA which pins memory
- Allocate global memory
- Global memory has smaller regions for specific uses:
  - **Constant memory:** If we put variables here they will be cached for broadcast in the SMs' uniform caches
  - **Local memory:** Any of a thread's local variables that can't go in registers go here

## Threads, Blocks and warps overview

- **Threads** are organised into **blocks**
- **Blocks** are organised into **Grids**
- The **thread** in a **block** and/or the **blocks** in a **grid** can be organised in 1D, 2D or 3D structures
- Whole **blocks** are allocated to SMs as they become available
- When a **block** is allocated to a SM, it is divided into **warps** of 32 **threads**
- We want to avoid the situation where all **warps** in a SM are waiting and not runnable
- Two strategies:
  1. Maximise the number of **warps** in an SM (*its occupancy*) by optimising block size
  2. Reduce the time spent waiting for memory access by using *coalescing* and *shared memory*
- A **block** consists of multiple **threads**
- A **grid** consists of multiple **blocks**

## Optimisation techniques:

- Increased thread **occupancy**
- **Coalesced** memory access
- **Shared** memory

## 2D blocks and grids

- If you use 32x32 blocks of 1024 threads:
- How many blocks would you need for a 2D array of 2,450 x 3570?
- x direction:  $2,450 / 32 = 76.5 = 77$
- y direction:  $3,570 / 32 = 111.5 = 112$
- **Cannot have half blocks so round up**
- If you round up, you need an if statement to catch the threads that will be out of bounds
- 2D and 3D blocks and grids will all converted to 1D by CUDA automatically

## Block assignment

- Each SM has one or more **blocks** assigned to it ('resident') at any given time
- Only whole **blocks** can be assigned to SM

- Therefore all threads in a block will run on the same SM
- The number of blocks/threads that can be resident in an SM depends on the device's Compute Capability

## Warps

- When a SM is assigned a block of threads it breaks it down into warps
- Each **warp** is a set of 32 threads that will be executed in parallel
  - Threads are assigned to warps by dividing the linearised blocks into 32 thread chunks

## Warp Scheduling

- Understanding warp scheduling means you can improve performance
- At any given time, a warp can be in 3 states:
  - **Running:** Currently being executed by the SM
  - **Waiting:** Can't run because it is waiting on something, usually memory access
  - **Ready:** Not running but not waiting on something
- Ideal situation is when one is running, there are two more ready and more waiting
- Worst situation is to have all SMs waiting, effectively meaning SM is idle

## Reducing stalls

- Most effective way to reduce stalls (assuming that it is waiting on memory)
- Is to use memory more efficiently
- Coalesced memory reads and shared memory
- But also important to maximise the number of warps in an SM
  - **its occupancy**
- The more warps there more chance there's one ready to be executed

## Calculating optimal occupancy

### Need:

- Max no of threads per block
- Max no of resident blocks per multiprocessor
- Max no of resident threads per multiprocessor

Following block sizes:

- 8x8
- 64 threads per block (2 warps)

- Each SM can only contain 8 blocks
- Max number of resident threads is  $8 * 64 = 512$ 
  - Much less than the capacity of 1536
- 16x16
- 256 threads per block (8 warps)
- Each SM can only contain 8 blocks
- Max number of resident threads is  $6 * 256 = 1536$ 
  - Perfect fit
- 32x32
- 1024 threads per block (32 warps)
- Max cap is 1536
- So cant have any more than one

### Coalesced memory access

- If threads in a warp simultaneously access memory addresses that are close together (128 bytes)
- Then the accesses are coalesced into one transaction
  - Much faster
- You have to plan how data is stored in memory so that the data that the threads need are next to each other

### Vector major order

- Each full vector is placed next to another from start to end element
- To add all first elements of the vectors to a scalar
- threads have to skip the length of the vector to get to the next first element
- Means wasted N number of loaded elements
- To access  $\text{globalId} * N + i$
- Global Id = thread
- N = size of vector
- i counter for each element in array

### Component major order

- Each first element is stored next to each other
- Then each second element, each third, etc.
- To add all elements by a scalar value
- Can access all consecutive memory addresses for the number of vectors to add
- To access:  $i * M + \text{globalId}$

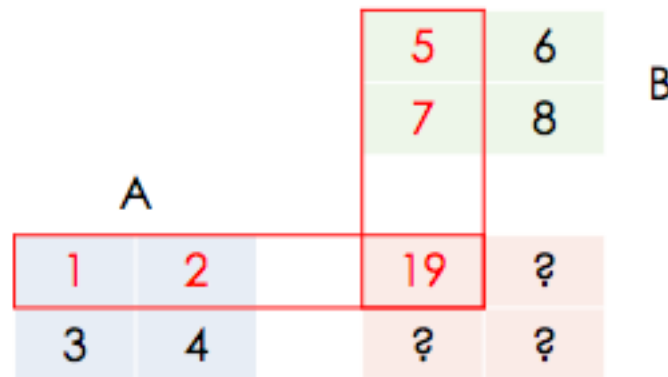
## Divergence

- Try to avoid branching statements in kernel code
- Only occurs when threads in the same **warp**
- All threads in the same warp share the same instruction pointer
- Therefore, all threads have to be on the same instruction
  - Branching causes threads to take different paths through the code depending on their value
- This can make it very inefficient
- The code should run serially at all times possible
- Try to change so only certain blocks access code meaning that all other warps in other blocks don't have to
- If threads do have to wait for other threads after divergence
  - They're marked as de-active
  - until they are all at the same point

## Syncthreads

- This causes threads to pause and wait until this point
- Waits until all threads are together
- Resumes code
- Can be dangerous
  - Don't write it in a branch as some threads may never reach it

## Matrix multiplication




- for first element in top row, take first row and first column and get the dot product
- for second element in top row, take first row and second column
- Repeat for second row

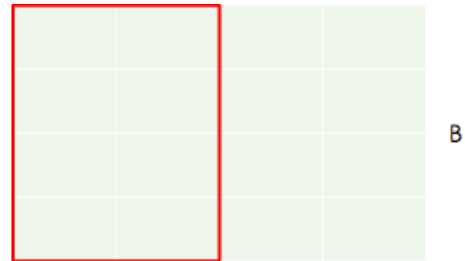
## Shared memory matrix multiplication

- Each row of the block
- The threads need to access the entire row of the matrix that they're multiplying
- Likewise with the columns
- Therefore, a whole block's threads need to access both columns and rows twice
- If you load first two rows and first two columns for first

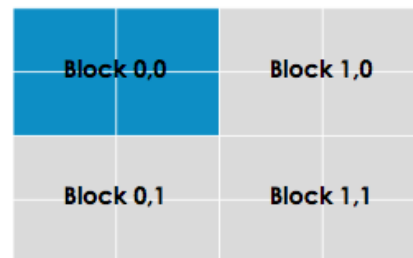
- These rows and columns are each accessed by this block twice in total.
- If we just load them once (into *shared memory*), global memory access is *halved*.



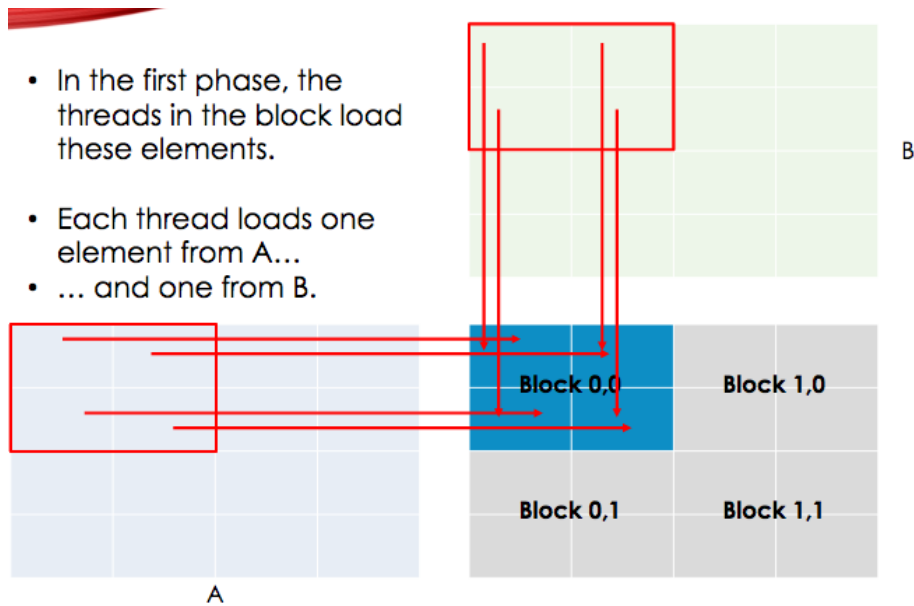
A



B







- In the first phase, the threads load the size of their block that they're calculating from each one of the matrix
- It calculates half its dot product
- If the matrices were 6x6 then  $n = 6$  and the blocks were 2x2
- we would need 3 phases
  - Phases are the turns in which the thread loads its elements into the block from the memory
- num of phases =  $n/m$

Compute to global memory access (CGMA)

$$CGMA = \frac{\# \text{ floating point operations}}{\# \text{ global memory accesses}}$$

- How many floating point operations will a block do?
- How many global memory accesses will it do

E.g: for a convolution mase

- $Db = 4 = \text{No. of threads in a block}$
- $M = 5 = \text{Mask size}$

|2|4|**5**|**2**|**3**|**7**|**2**|9|1

- In bold is the no of threads in the block =  $D_b$

**How many computations?**

- Each thread does  $M$  multiplications
  - one for each of the mask
- And then does 4 additions =  $M - 1$
- There are  $D_b$  threads in a block
  - So:

$$D_b * (M(M-1)) = D_b * 2M - 1$$

**How many global memory accesses?**

- Each thread loads  $M$  elements
- Also has to store 1 element (the result)
- $D_b$  threads in a block
  - So:

$$D_b * (M + 1)$$

$$CGMA = \frac{D_B(2M-1)}{D_B(M+1)} = \frac{2M-1}{M+1} \approx 2$$

Rubbish!

**Using shared memory:**

- Only the global access memory part of the ratio changes
- Each block as a whole loads  $D_b + M - 1$
- Has to store  $D_b$  (its results)
  - So:

$$2D_b + M - 1$$

$$CGMA = \frac{D_B(2M-1)}{2D_B+M-1} = \frac{2M-1}{2+\frac{M-1}{D_B}} \approx \frac{2M}{\frac{M}{D_B}} = \frac{2MD_B}{M} = 2D_B$$

**Task parallelism**

- A task is split up into different sections where different workers work on different parts of the same task at the same time

### Data parallelism

- A task is processed from start to finish by one thread/worker
- But potentially different data is used for each task

### Single instruction stream, single data stream (SISD)

- The CPU does **one operation** at a time on **one piece of data**

### Single instruction stream, multiple data stream (SIMD)

- **Multiple cores** that process **multiple data streams** in parallel
- **BUT** each core executes the same instruction at once

### Multiple instruction stream, single data stream (MIMD)

- The **same data** is read from memory into **multiple processing units** each of which have their own instructions

### Multiple instruction stream, multiple data stream (MIMD)

- Each processing unit can **access different data from memory** and run **different instructions**

### Shared memory system

- Each processing core can **directly access the same memory**

### Distributed memory system

- Each processing core has its **own private memory**
- It can only access **another's indirectly**

### Hybrid memory system

- **Groups** of processing cores **share memory**
- **But** memory is **not** shared between groups

### Speed-up

- How long does a parallel version of a program take to run vs a serial version
- $T_S$  - how long the **serial** program takes to run.
- $T_P$  - how long the **parallel** program takes to run.

$$S = \frac{T_S}{T_P}$$

### Linear speed-up

- If parts of the program can be done completely independently
- Adding more processors means that the task becomes faster
- Good for scaling as you can just add more processors

$$T_P = \frac{T_S}{p} \quad \text{so} \quad \text{Speedup} = S = \frac{T_S}{\frac{T_S}{p}} = p$$

### Parallel efficiency

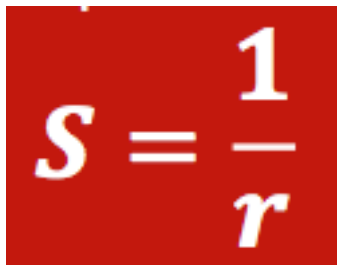
- As the processor count increases
  - So do the associated overheads
- Although speed up increases
  - It gets less so with each processor added
- Measure it with efficiency
- For linear speed-up,  $E = 1$
- For sublinear speed-up,  $E < 1$

$$E = \frac{\text{Speedup}}{\# \text{ Processors}} = \frac{S}{p}$$

## Amdahl's Law

*‘There is a limit to how fast you can do something even if you add more processors’*

- It is limited by the time spent processing the section of code that can't be serialised
- For the same problem size, as the number of processors increases speed-up is limited to:
  - *where  $r$  is the proportion of code that cannot be paralised*

A red rectangular box containing the white text of the formula  $S = \frac{1}{r}$ . The 'S' is large and bold, followed by an equals sign, then a fraction with '1' in the numerator and 'r' in the denominator.

- No matter how many processors we have the code will always be limited to:

**1 / proportion of code that can't be paralised**

- It assumes you want to process the same amount of data as quickly as possible by adding processors
  - *task-parallel approach*

## Gustafson's Law

*‘By adding more processors you can always process more data in a given period of time’*

- Assume that the problem size (amount of data to be processed) increases with the number of processors ( $p$ )
- Divide parallel program execution time into two parts:

$$T_P = a + b$$

Time taken in serial execution      Time taken in parallel execution

**Example:**

- CUDA program that runs in parallel across 1,000 cores, processing 1,000 pieces of data
- Starting up CUDA, copying data to/from GPU, launching kernel takes 100ms
  - $a = 100\text{ms}$
- The kernel runs (on all cores in parallel) for 500ms
  - $b = 500\text{ms}$

$$T_P = 100 + 500 = 600\text{ms}$$

- What is  $T_S$ ? - Whatever the cores are doing to the data takes 500ms - So if we only have one core, each piece of data has to be processed sequentially:

$$T_S = 100 + 1000 \times 500 = 500100\text{ms}$$

- If the problem size scales with the processor count ( $p$ )
- The speed-up is limited to:
  - where  $r$  is the proportion of the program that can't be parallelised

$$S = p + r - rp$$

**Message Passing Interface MPI**

- MPI is an application programming interface (API) for *distributed memory* parallel programming
- Usually MPI follows a *single program multiple data (SPMD)* approach
  - Write one program and spawn multiple copies of it:

```
mpiexec -n 100 program.exe
```

- But can also be used for a *multiple program multiple data (MPMD)* approach:

```
mpiexec -n 50 program.exe : -n 50 otherProgram.exe
```

## MPI terminology

- **Communicator:** A group of processes that can talk to each other
  - Default: `MPI_COMM_WORLD` - contains all processes
- **Rank:** An integer uniquely identifying a process with its communicator
- **Tag:** A (user-defined) integer attached to a message that can be used to indicate the type of message

## Threads v Processes

- Every individual program is a separate *process*
- Processes cannot normally access each others' memory
  - *distributed memory*
- Communicate in various ways
- Each process has one main *thread* of execution
  - But has the ability to spawn additional ones
- Each thread can do something different
  - *run different bits of code*
- All threads can access their processes' memory
  - *shared memory*

## Threads

- 'lightweight'
- Low memory overhead
- Can switch between them quickly
- Shared memory means data can be exchanged very easily and quickly
- All need to be on the same machine

## Processes

- ‘heavyweight’
- Higher memory overhead
- Switching takes longer
- Distributed memory communication is harder and slower
- Can be running on different machines
- Communicate in various ways:
  - Network sockets
  - Files
  - Named pipes (OS feature)

## Blocking

- MPI\_Send and MPI\_Recv are both potentially *blocking* functions
- This means that every MPI\_Send on one process must have a matching MPI\_Recv on another process
  - **Or the process will hang**
- Messages are ‘*nonovertaking*’

## Connection topologies

- Connecting multiple processing cores that are not on the same PC

*Types of topologies:*

- **Fully connected:** Every PC is connected to one another
- **Ring:** Every PC is connected to two other PCs to make a circle or a ring
- **Thin tree:** One computer at the top, for every level down, the number of PCs doubles, one connection each
- **Fat tree:** Same layout as thin tree however every level you go up the connections double
- **Torus:** A doughnut shape where every side is connected

## Topology terminology

### Diameter:

- The maximum path length between a pair of nodes
- High values of diameter cause more latency

### Bisection width/Band width:



- The number of links you cut to divide the network into two equal halves
- To get the bisection bandwidth you multiply the bisection width with the bandwidth
- A high value is better as more links mean more resilience and high bandwidth helps with algorithm computation

#### Valency:


- How many connections each node makes
- Want to keep this to a reasonable amount
  - 6 or 7 max

#### Link count:

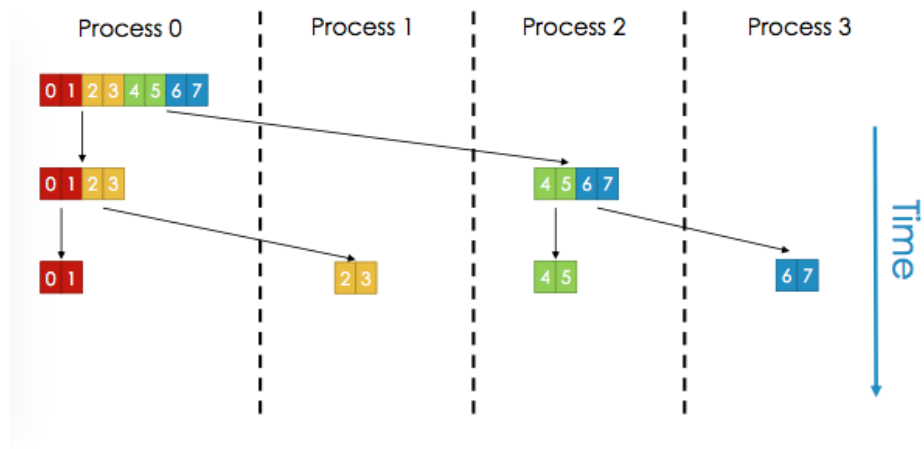
- How many connections the network has in total
- Shows how many wires are needed

#### Distributing data - binary tree

- Set  $M$  = the number of processes
- At every time step
  - Divide  $M$  by two
  - If a process with rank  $i$  has the data, it sends it to the process with rank  $i + M$
- Very efficient at transmitting data
  - takes  $\log_2 N$  time steps

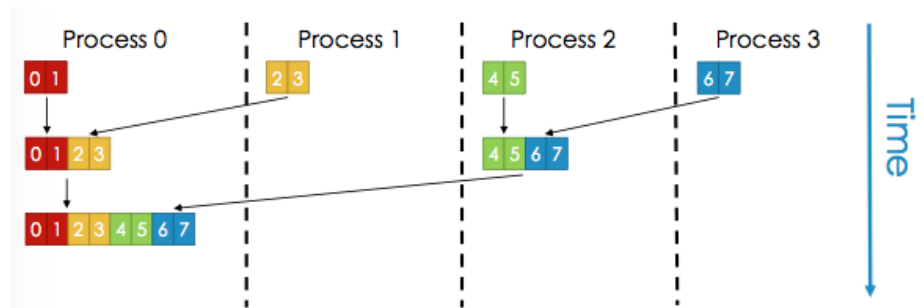


### MPI\_Scatter




### MPI\_Gather

- Does the opposite of MPI\_Scatter
- Collects data from arrays distributed across all the processes into one big array



### MPI\_Reduce

- Take an array of values and reduce it to a single value




- Divide (active) processes into pairs.
- 2<sup>nd</sup> process in pair gives its value to 1<sup>st</sup> process, goes inactive.
- 1<sup>st</sup> process in pair takes the bigger of the two values.
- Repeat.

## MPI\_AllReduce

Pairs of individual processes compare values...

Pairs of blocks of 2 processes compare values...

Pairs of blocks of 4 processes compare values...

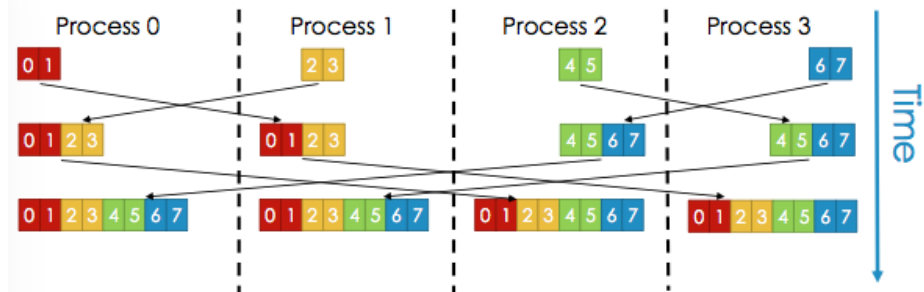


Requires  $\log_2 N$  steps – same as reduce or broadcast alone!

(But does require more data transfers...)

## MPI\_AllGather

- Collects data from arrays distributed across all the processes into one big array *that all processes have a copy of*.
- Could be implemented using a butterfly arrangement:



## Topics that came up v topics that haven't come up

*italic text means they've appeared twice*

Have appeared	Haven't appeared
<i>static allocation</i>	Gustafon's law
pinned memory	what a SM is
which types of memory	difference between host and device
instruction level parallelism	difference between static and dynamic
constant memory	disadvantages of performance with branching
how many threads will be created for a kernel launch	how threads can cooperate to load data into shared mem
data organisation into cache	why CGMA is important
<i>row/component major order</i>	convolution
maximising occupancy	how shared memory can be optimised calculating a convolution

Have appeared	Haven't appeared
—	—
maximum number of blocks that can be resident in multiprocessor	difference between task and data parallelism
—	—
<i>coalesced memory access</i>	sublinear & super linear
—	—
<i>CGMA</i>	difference between threads and processes
—	—
synctreads	<code>MPI_Send</code> & <code>MPI_Recv</code>
—	—
<i>divergence</i>	how MPI functions are blocking and non-overtaking
—	—
network topologies and their properties	different MPI collective communication functions
—	—
register spilling and effect on other memory because of it	why and how collective communication functions can be efficient
—	—
CGMA of code	—
—	—
SISD,SIMD,MISD,MIMD	—
—	—
speedup	—
—	—
efficiency	—
—	—
amdahls	—
—	—
linear speedup	—
—	—
shared memory	—
—	—
general formulas for topologies	—