

# SOFT354 - CUDA memory

Date: 06-10-16

## Memory matters

- Compute to global memory access ration
  - How many floating point operations does a program do for every global memory access operation
- GPUs have > 200GB/s global memory bandwidth
  - 50,000 floating point values per second
- Copying from CPU to the GPU is even slower than global memory access
- We need to minimise access to global memory

## Host and device memory

- **Host memory** is the PC's normal RAM
- **Device memory** is the “video” RAM on the GPU
- Code in a *kernel* can **only** access **device memory**
- Host code (main() function) can access device memory using functions like cudaMemcpy

## Sin(x) example

- In the workshop we wrote a program to compute sin(x) for a lots of values of x
- each **thread** computes one value
- two pairs of arrays:
  - One on the host: **input, result**
  - One on the device: **d\_input, d\_result**

Steps: 1. Serially initialise input data on host 2. Copy input to GPU - Using cudaMemcpy 3. Kernel runs, computes sin(x) in parallel - On each thread 4. Copy result back to host to work with

## Static vs Dynamic allocation

- Arrays can be *statically* or *dynamically* allocated
- With **static allocation** the size of the array must be known at compile time
  - Space is reserved in the program's memory map
  - `float staticArray[10];`
- With **Dynamic allocation** the size of the array can be calculated at runtime

- `float* dynamicArray = (float*)malloc(sizeof(float));`
- how many bytes you want
- **Array is just a pointer pointers are just integers**
- In both cases the variable is just the memory address of the first element in the array

	Allocate in <u>Host Memory</u>	Allocate in <u>Device Memory</u>
Static	<code>float h_array[10];</code>	<code>__device__ float d_array[10];</code>
Dynamic	<code>float* h_array = (float*)malloc(10*sizeof(float));</code>	<code>float* d_array; cudaMalloc(&amp;d_array, 10*sizeof(float));</code>

- `cudaMalloc` returns a `cudaError_t` - Not the address of the allocated memory
- So you need to pass a pointer to a pointer - If you use `malloc` inside a kernel - It will allocate memory on the device

### Allocating memory in kernels is rare

- A common mistake is put the `cudaMalloc` call inside the kernel
- This is possible but not desired
- Every thread would allocate enough space for the whole array
  - would run out of space
- Instead we allocate enough space to hold one copy of the array in device memory
- and each thread accesses a different bit of it

### Freeing dynamic memory

- Dynamically allocated memory isn't automatically cleaned up
- `free(array)`
- `cudaFree(array)`

### Copying to static device arrays


- When you use a static allocation for device memory
- the variable isn't a pointer to a memory address on the GPU, it's a symbol
- This means you can't directly pass the variable into `cudaMemcpy`
  - Use `cudaGetSymbolAddress`
  - `cudaMemcpyToSymbol/cudaMemcpyFromSymbol`

## Accessing static memory from a kernel

- if device memory was allocated statically in the host code
- It allows you access it *globally* on the device
  - i.e. not have to pass into functions
- **Better practise**

## Pinned memory

- Technique for speeding up RAM < - > memory transfers
- Transfers between the RAM and GPU can be very fast
  - They use direct memory access to do the copy without involving the CPU
- **but**, operating systems used **paged virtual memory**
- The address you get from malloc() doesn't correspond to a physical address in RAM
  - but to a “page” that can be moved around
  - Which is no good for a DMA transfer
- So when transferring from RAM to GPU
  - Cuda first copies the data into **pinned** RAM
    - \* Where it doesn't get moved



- When you allocate array on host - Can specify that it should be pinned - When you do any type of transfer - It can use DMA - Very fast - like malloc, dynamically allocates an array in RAM - Unlike malloc, the memory will be pinned - **Use cudaFreeHost to deallocate when the memory isn't needed anymore**

## Other CUDA memories

- So far we've only allocated in two ways

- **Static**
  - **Dynamic**
- Both of these approaches allocate memory in the GPU's **global memory (RAM)** which is:
  - By far the **biggest** area of memory
  - Also the **slowest**
- There are other **smaller, faster** memories that we can use to dramatically increase performance